

## Introduction

The top-level instructions include information for the installation and functional validation of OpenMPI, with installation on the head and compute nodes.

```
$> dnf -y install openmpi openmpi-devel gcc-c++  
$> dnf -y --installroot=$CHROOT install openmpi openmpi-devel gcc-c++  
$> packimage centos8-x86_64-netboot-compute  
$> pdsh -w c[1-2] reboot
```

Successful installation is confirmed with functional validation testing. There were basic connectivity tests by pinging to ensure the compute nodes can reach the head node, and using a sample MPI script to demonstrate that the compute nodes can communicate with each other. In this document we will move beyond functional testing and validate network performance to ensure it meets expectations. Here, performance validation is used to determine the speed of network communications over a given interface. In other words, this document will validate that the communication network(s) are performing at an acceptable level. Performance will be measured using the using the OSU Micro-Benchmark suite, executed with the open-source OpenMPI, in addition to the IMB benchmarks from Intel OneAPI.

## Ethernet: OpenMPI

This section will validate the performance of the standard Ethernet connection using OpenMPI. It assumes a head node connected to two compute nodes - c1 and c2 - in a minimal cluster setup using the instructions in the top-level cluster\_setup documentation, with no other connections on the cluster such as Infiniband or high-speed Ethernet connection. Using root is not recommended unless otherwise noted.

Install the OSU Micro-Benchmarks that will be used for performance testing. The following commands install the OSU benchmarks in the user's home directory.

```
$> wget https://mvapich.cse.ohio-state.edu/download/mvapich/osu-micro-benchmarks-5.8.tgz  
$> tar -xzf osu-micro-benchmarks-5.8.tgz  
$> cd osu-micro-benchmarks-5.8/  
$> ./configure CC=/usr/lib64/openmpi/bin/mpicc CXX=/usr/lib64/openmpi/bin/mpicxx  
$> make  
$> make install exec_prefix=~/.osu_benchmarks_openmpi
```

We will validate performance by using one-sided (RMA) communication for lower overhead. The `osu_put_bw` test is adequate for this purpose. Use the following sample script to execute the `osu_put_bw` test on two compute nodes:

```
#!/bin/bash -l
#SBATCH -N 2
#SBATCH -J perf_test
#SBATCH -p normal
#SBATCH -t 20
#SBATCH -o osu_perf_test.out
#SBATCH -e osu_perf_test.err

export PATH=/usr/lib64/openmpi/bin:$PATH
mpirun -n 2 -N 1 -mca btl self,tcp
~/osu_benchmarks/libexec/osu-micro-benchmarks/mpi/one-sided/osu_put_bw
```

Submit the script on the head node with sbatch. A sample output from the osu\_put\_bw benchmark is included below:

```
# OSU MPI_Put Bandwidth Test v5.8
# Window creation: MPI_Win_allocate
# Synchronization: MPI_Win_flush
# Size      Bandwidth (MB/s)
1           0.20
2           0.40
4           0.81
8           1.58
16          2.89
32          5.67
64          10.22
128         21.14
256         40.65
512         65.06
1024        84.70
2048        98.46
4096        106.42
8192        110.92
16384       113.80
32768       115.32
65536       116.06
131072      116.53
262144      116.76
524288      116.88
1048576     116.92
2097152     116.95
4194304     116.96
```

The output from the benchmark is listed in MB/s. For easier comparison we will convert to Gb/s. Using the highest listed bandwidth output, convert as followed:

$$\frac{116.96 \text{ MB}}{1 \text{ s}} \times \frac{1 \text{ GB}}{1000 \text{ MB}} \times \frac{8 \text{ Gb}}{1 \text{ GB}} = \frac{.9357 \text{ Gb}}{1 \text{ s}}$$

The calculated .93Gb/s is close to the theoretical peak rate of 1Gb/s for the ethernet connection.

## Ethernet: Intel MPI

The previous section validates performance of the default Ethernet connection using OpenMPI. Alternatively, the Intel OneAPI Toolkit can be used for the same purposes. Performance validation with the OneAPI Toolkit uses Intel MPI with IMB (Intel MPI Benchmarks), instead of openMPI with OSU benchmarks.

Install Intel MPI from the OneAPI Toolkit:

```
$> dnf config-manager --add-repo https://yum.repos.intel.com/oneapi
$> rpm --import https://yum.repos.intel.com/intel-gpg-keys/GPG-PUB-KEY-INTEL-SW-PRODUCTS.PUB
$> dnf install intel-oneapi-mpi-devel
```

The default installation for the MPI executables will be `/opt/intel/oneapi/mpi/latest/bin`, with the IMB benchmarks located at `/opt/intel/oneapi/mpi/latest/benchmarks/imb`. To save space on the diskless compute nodes, the Intel MPI folders on the head node will be shared with the compute nodes through NFS instead of being installed into the compute image.

```
$> echo "10.10.1.10:/opt/intel /opt/intel nfs nfsvers=3,nodev,nosuid 0 0" >> $CHROOT/etc/fstab
$> echo "/opt/intel *(ro,no_subtree_check,fsid=13)" >> /etc/exports
$> systemctl restart nfs-server
$> packimage centos8-x86_64-netboot-compute
$> pdsh -w clx[1-2] reboot
```

Note the `fsid=13`. This number may need to be changed, depending on other folders shared. Check `/etc/exports` to see if FSID 13 has been reserved for a different folder. If it is, then change to the lowest number that is not being used.

Use the following sample script to execute one of the IMB RMA tests - PingPong - on two compute nodes:

```
#!/bin/bash -l
#SBATCH -N 2
#SBATCH -J perf_test
#SBATCH -p normal
#SBATCH -t 20
#SBATCH -o imb_perf_test.out
#SBATCH -e imb_perf_test.err

source /opt/intel/oneapi/mpi/latest/env/vars.sh
mpirun -np 2 -ppn 1 IMB-P2P PingPong
```

After submitting the script with `sbatch`, a sample output is shown below:

```
#-----
# Benchmarking PingPong
# #processes = 2
#-----
#bytes #repetitions t[usec] Mbytes/sec Msg/sec
0 100000 49.81 0.00 20077
1 100000 48.05 0.02 20813
2 100000 48.35 0.04 20681
4 100000 54.19 0.07 18454
8 100000 56.44 0.14 17717
16 100000 59.67 0.27 16759
32 100000 60.38 0.53 16561
64 100000 72.15 0.89 13860
128 100000 88.67 1.44 11278
256 100000 63.22 4.05 15818
512 100000 62.63 8.18 15967
1024 100000 62.91 16.28 15895
2048 100000 69.48 29.48 14393
4096 100000 126.80 32.30 7887
8192 100000 189.69 43.19 5272
16384 51200 231.20 70.87 4325
32768 25600 388.93 84.25 2571
65536 12800 625.94 104.70 1598
131072 6400 1205.68 108.71 829
262144 3200 2457.57 106.67 407
524288 1600 4687.17 111.86 213
1048576 800 9232.21 113.58 108
2097152 400 18220.72 115.10 55
4194304 200 36276.94 115.62 28
```

Similar to the OSU benchmarks, IMB output uses Mb/s, which we can convert:

$$\frac{115.62 \text{ MD}}{1 \text{ s}} \times \frac{1 \text{ GB}}{1000 \text{ MB}} \times \frac{8 \text{ Gb}}{1 \text{ GB}} = \frac{.925 \text{ Gb}}{1 \text{ s}}$$

.925Gb/s is close to the theoretical peak rate of 1Gb/s for the ethernet connection.